

Comparative Analysis of a Tagged Interval Tree Implementation for Text Editing

- **Introduction: Interval Trees and the Necessity of Tagging in Text Editing**

Interval trees are a specialized tree data structure primarily designed to efficiently manage and query intervals, which are defined by a start and an end point.¹ Their fundamental operation involves finding all intervals that overlap with a given query interval or a specific point in time or space. This capability makes them particularly useful in a variety of applications, including windowing queries in geographical information systems, where one might need to find all roads within a specific map viewport, or in three-dimensional graphics, to identify all visible elements within a scene.² The core strength of interval trees lies in their ability to perform these overlap searches significantly faster than a naive linear scan of all intervals.

The provided code introduces a specific type of interval tree tailored for tagging text. This implementation emphasizes several key properties outlined in its algorithm description: Non-overlapping, Sorted, Merged, Inside, and Tag. These properties are crucial for maintaining a consistent and efficient representation of text where specific segments are associated with tags indicating formatting or semantic meaning. The enforcement of these strict properties suggests a design philosophy aimed at creating a canonical representation of the tagged text. For instance, the "Non-overlapping" and "Merged" properties likely aim to prevent ambiguities and redundancies by ensuring that tagged regions at the same level of the tree are distinct and that contiguous regions with the same tag are unified. The "Sorted" property likely facilitates more efficient traversal and searching within the tree structure.

In the context of text editing, tagging plays a vital role in applying rich formatting (such as bold, italics, or underline) or in embedding semantic information within specific ranges of text.⁵ For a text editor to be responsive and efficient, especially with large documents, the underlying data structure must allow for quick addition, removal, and querying of these tags. Operations should ideally not require processing the entire document for every change or query. Interval trees, with their inherent ability to efficiently handle interval-based data, offer a promising approach to meet these performance requirements. By organizing the tagged text as a set of intervals within an interval tree, the editor can leverage the tree's structure to quickly locate and modify tagged regions based on their start and end positions.

This report aims to provide a detailed comparative analysis of the provided tagged interval tree implementation. It will examine its core structure, the

algorithms it employs for managing intervals and tags, and the specific properties it enforces. By comparing these aspects with the established concepts of interval trees and the features of existing JavaScript interval tree libraries, this analysis will identify the similarities and unique characteristics of the provided implementation, offering insights into its design and potential suitability for text editing applications.

- **In-Depth Analysis of the Provided Tagged Interval Tree Implementation**

The provided code defines two primary classes: `IntervalNode` and `TaggedIntervalTree`. The `IntervalNode` class serves as the fundamental building block of the tree. Each instance of this class represents a segment of the text and holds three key pieces of information: the interval, defined by a start and an end value, indicating the range of text it covers; an optional tag, which is a string associated with this segment, signifying the formatting or semantic information; and an array of children, which are other `IntervalNode` instances. This children array is particularly significant as it suggests a hierarchical structure within the tagged text representation. The nesting of children within a parent node implies that tags can be applied to segments of text that are already within a tagged region, mirroring the common practice in rich text editors where, for example, a word can be bolded within an italicized phrase. The interval of each child node is expected to be entirely contained within the interval of its parent, reflecting the "Inside property" mentioned in the algorithm description.

The `TaggedIntervalTree` class encapsulates the overall tree structure. It contains a single property, `root`, which is an instance of `IntervalNode` representing the root of the entire tree. The constructor initializes this root node with a start and end value, typically representing the boundaries of the entire text being managed. The algorithm description for this class explicitly mentions five key properties that the tree is designed to maintain: "Non-overlapping property," stating that intervals at the same level of the tree should not overlap; "Sorted property," indicating that intervals at the same level must be sorted based on their start values; "Merged property," suggesting that contiguous intervals at the same level with the same tag should be merged into a single interval; "Inside property," as discussed earlier, requiring that a child's interval is always within its parent's interval; and "Tag property," specifying that a node cannot have the same tag as its parent. These properties collectively define a structured and consistent way of representing tagged text within the interval tree.

The `addTag` method in the `TaggedIntervalTree` class is responsible for applying a given tag to a specified interval of text. It begins by performing a validation check to ensure that the provided interval (start and end points) is valid, returning if the start is greater than or equal to the end. The core logic of adding a tag is

delegated to the `_addTagDFS` method, which employs a Depth-First Search (DFS) strategy to traverse the tree and find the appropriate location to insert the new tag. The `_addTagDFS` method first ensures that the operation stays within the bounds of the current node's interval. If the resulting interval is invalid (start greater than or equal to end), the method returns. It then checks if the current node already has the same tag as the one being added. If so, no further action is needed for this branch, adhering to the principle of avoiding redundant tagging at the same level. If the current node has no children, a new `IntervalNode` is created with the specified tag and interval and is added as a child to the current node. When the current node already has children, the algorithm needs to determine the correct insertion points for the new tagged interval. It iterates through the existing children, keeping track of the current position within the overall interval. Before the first child, if there's a gap between the current position and the start of the first child's interval, a new insertion point is identified. Similarly, between each child, if there's a gap, another insertion point is recorded. If the current position overlaps with a child's interval, the `_addTagDFS` method is called recursively on that child to handle tagging within that subtree. After processing all existing children, if there's still a portion of the target interval remaining after the last child, a final insertion point is noted. Once all potential insertion points are identified, the algorithm processes them in reverse order. For each insertion point, a new `IntervalNode` with the given tag and interval is created. Before inserting this new node, the algorithm attempts to merge it with any existing neighboring children that have the same tag. If the previous sibling has the same tag and its end point is at or after the start of the new node, the previous sibling's interval is extended to cover the new node's interval, and if the next sibling also has the same tag and overlaps, it is also merged into the previous sibling, along with its children. If no merge is possible with neighbors, the new node is inserted at the calculated index in the current node's children array. Finally, after potentially adding new children, the `_processChildrenWithSameTag` helper function is called. This function iterates through the children and checks if any child has the same tag as the current node. If it does, those children are removed, and their grandchildren (that do not have the same tag as the current node) are moved up to become direct children of the current node. This step directly enforces the "Tag property" by preventing a node from having a child with the same tag.

The `removeTag` method is designed to remove a specific tag from a given interval. Similar to `addTag`, it starts with a check for invalid input intervals. The core removal logic resides in the `_removeTagDFS` method, which also uses recursion to traverse the tree. This method takes the current node, the tag to remove, and the

start and end of the removal interval as input. It first adjusts the removal interval to be within the bounds of the current node's interval. If there is no overlap between the adjusted interval and the node's interval, it returns indicating no removal occurred in this subtree. If the current node's tag matches the tag to be removed, the method proceeds with different cases based on the relationship between the removal interval and the current node's interval.

If the removal interval is entirely within the current tagged interval (not touching the start or end), the algorithm iterates through the current node's children, categorizing them into those before the removal interval, those inside it, and those after it. Children that overlap with the removal interval are processed recursively. For those children that are completely removed or split due to the recursive call, their "rehook nodes" (explained below) are collected. A new IntervalNode is created for the portion of the current tag's interval before the removal interval, and another for the portion after. The children that were before the removal interval are assigned to the "before" tag node, and those after to the "after" tag node. The children that were inside the removal interval (and not entirely removed) are also collected. All these newly created and existing children are added to a list called rehookNodeList, sorted by their start position, and then merged if possible. The method returns a state indicating that the removal interval was inside, along with the rehookNodeList and an empty remaining interval, signaling the termination of DFS for this branch. The parent node, upon receiving this state, will remove the current tag node and add the nodes from the rehookNodeList as its children.

If the removal interval starts at or before the current tag's start and ends within it, the current node's interval is adjusted to start at the end of the removal interval. Affected children (those entirely within or partially overlapping the removal interval) are processed recursively. If a child is entirely removed, its rehook nodes are collected and added back to the current node's children if their interval is after the removal interval. The method returns a state indicating a left-side removal and the remaining interval after the removal. The parent node will add any returned rehook nodes and continue DFS with the remaining interval. A similar logic applies if the removal interval starts within the current tag and extends to or beyond its end. The current node's interval is adjusted to end at the start of the removal interval, affected children are processed, and a state indicating a right-side removal is returned along with the remaining interval before the removal.

If the removal interval completely covers the current tag's interval, the algorithm processes all children. If a child overlaps with the removal interval, it is also subjected to tag removal recursively. Any rehook nodes returned from these

recursive calls are collected. The method returns a state indicating that the entire node should be removed, along with the collected rehook nodes. The parent node will then remove this node and add the rehook nodes as its own children. In cases where the current node does not have the tag to remove, the algorithm iterates through its children and recursively calls `_removeTagDFS` on each child that overlaps with the removal interval. If a child's tag is removed, the parent node updates its children list, potentially replacing the removed child with the rehook nodes it returned. The process continues with any remaining portion of the removal interval. The justification for the non-overlapping property after removal is that since all children's intervals are initially inside their parent's non-overlapping intervals, re-hooking them to the parent of the parent maintains this non-overlapping characteristic at that higher level.

The `hasTag` method checks if a given interval has a specific tag. It uses a recursive `_checkTagDFS` method that traverses the tree. It returns true if a node with the specified tag is found that fully contains the given interval. The search continues in the children of the current node if they overlap with the given interval. The `getFormattedText` method is responsible for generating a formatted text string from the tagged interval tree. It uses a helper function `collectMarkers` that performs a traversal of the tree and records the start and end positions of each tagged interval along with the tag itself, creating "markers." These markers are then sorted by their position, with closing tags prioritized at the same position. Finally, the method iterates through the sorted markers, constructing the output string by adding the text between markers and inserting the appropriate opening and closing HTML-like tags.

- **Survey of Interval Tree Implementations and Concepts in Research Material**

Standard interval tree concepts revolve around augmenting a self-balancing Binary Search Tree (BST), such as a Red-Black Tree or an AVL Tree, to efficiently manage intervals.¹ A common augmentation involves storing the maximum high value (max) in the subtree rooted at each node. This additional information allows for optimized overlap queries, typically achieving a time complexity of $O(\log n)$ for operations like adding, removing, and searching for overlapping intervals.¹ The fundamental principle is to use the low value (start point) of an interval as the key for ordering within the BST.¹

Several JavaScript libraries provide interval tree implementations.

`node-interval-tree`¹² is one such library, built as an augmented AVL tree. It differs from the provided code in that it allows the insertion of the same interval multiple times, as long as the associated data is different. Each node in this tree maintains a list of records, where each record contains an interval and its associated data. It offers insert, search (for overlapping intervals), and remove functionalities. This

suggests its utility in scenarios where multiple independent pieces of information might be associated with the same text range.

Another prominent library is `@flatten-js/interval-tree`¹³, which is based on the algorithm described in "Introduction to Algorithms" by Cormen et al..¹⁴ This library focuses on storing key-value pairs, where the key is the interval. It supports standard operations like insert, remove, and search, and also provides methods for traversing and manipulating the tree, such as `forEach`, `map`, and `iterate`.¹⁵ This implementation appears to be a more general-purpose interval tree, adhering closely to the theoretical foundations without specific built-in support for tagging.

`interval-tree-type`¹⁶ is another JavaScript library that implements a red-black augmented interval tree. It offers features like support for unbounded intervals (intervals that extend to infinity) and custom value equality, which could be beneficial if tags were represented by complex objects or if there was a need to represent tags spanning the entire document. It includes a `queryWithinInterval` method for finding intervals within a given range. Other available JavaScript interval tree libraries include `rb-interval-tree`¹⁷, a balanced red-black tree; `static-interval-tree`¹⁸, a simple library for finding overlaps in a static set of intervals; and `davidisaaclee/interval-tree`¹⁹, which provides an immutable functional interface.

Segment trees are related tree-based data structures that are often confused with interval trees.¹ However, their primary focus differs. While interval trees are designed to store and query intervals based on overlap, segment trees partition the underlying space into segments, typically determined by the endpoints of the intervals, and store information relevant to these segments.²⁰ Segment trees are particularly efficient for range queries, such as finding the sum or minimum value within a given range.²¹ The fundamental difference lies in how they handle intervals: interval trees store the intervals themselves, whereas segment trees operate on the segments defined by the interval endpoints.¹ This makes interval trees a more direct fit for managing and querying tagged text, where the tags are associated with specific intervals of the text.

Interval trees have found applications in diverse domains beyond text editing. They are used in genomic interval processing to find overlapping gene regions or DNA sequences.¹⁰ They are also employed in windowing queries for maps and other graphical applications.² Scheduling applications can benefit from interval trees to manage time slots and detect overlaps.²³ Notably, `snippet`²⁷ specifically mentions the use of interval trees in an HTML rendering engine and editor to enhance performance in redrawing and minimizing memory usage. `Snippets`²⁸ and `28`²⁸ discuss the challenges of using a single string buffer for text editors and

the need for efficient insertion and deletion, areas where interval trees can provide advantages. The breadth of these applications highlights the general utility of interval trees in managing data associated with intervals.

- **Comparative Analysis: Unpacking the Similarities**

At a fundamental level, the provided tagged interval tree implementation shares the core concept of using a tree-based structure to manage intervals with other interval tree implementations.¹ The `IntervalNode` class, with its `interval` property storing the start and end points, is analogous to the node structures found in libraries like `node-interval-tree`¹² and `@flatten-js/interval-tree`¹³, which also store the boundaries of the intervals they manage. However, a key difference emerges in how associated data is handled. The provided code uses a direct `tag` property within the `IntervalNode`, indicating a specific focus on tagging, whereas other libraries like `node-interval-tree` use a more general `data` field to associate arbitrary information with an interval, and `@flatten-js/interval-tree` uses a value associated with an interval key.

The `addTag` algorithm in the provided code bears resemblance to the standard interval insertion process found in libraries such as `@flatten-js/interval-tree`.¹³

Both approaches involve traversing the tree to find the correct position for the new interval and potentially restructuring the tree to maintain its properties, such as sorted order. The provided `addTag` method further incorporates specific logic for merging contiguous intervals that have the same tag, a feature driven by the "Merged property" requirement. In contrast, the `removeTag` algorithm in the provided code is significantly more complex than a simple interval removal operation as seen in libraries like `node-interval-tree`.¹² The complexity arises from the need to handle various scenarios based on how the removal interval interacts with the tagged intervals, as well as the requirement to maintain the tree's structural integrity and the defined properties, including the "Inside" and "Non-overlapping" properties, through mechanisms like re-hooking child nodes. A distinctive aspect of the provided code is its direct handling of tags and the enforcement of the "Tag property," which restricts a node from having the same tag as its parent. This is a specific constraint tailored for text tagging that is not typically found as a built-in feature in general-purpose interval tree libraries. Libraries like `node-interval-tree` offer the flexibility to associate any type of data with an interval, which could include a tag, but they do not inherently enforce hierarchical constraints on this data. This suggests that the provided code adopts a more opinionated approach, embedding the tagging concept deeply into the tree structure and enforcing a particular model of tagged text.

The provided code places a strong emphasis on maintaining the "Non-overlapping property" at the same level of the tree through its merging

logic in the addTag method and the careful handling of children during tag removal. This contrasts with the primary function of many general-purpose interval tree libraries, which is to efficiently find overlapping intervals, implying that overlaps are expected and handled differently. While some libraries might offer functionalities for merging or splitting overlapping intervals ²⁵, the provided code's explicit enforcement of non-overlapping intervals at the same level appears to be a design choice aimed at simplifying the representation and potentially the rendering of tagged text, where overlapping tags at the same level might be considered redundant or invalid.

Finally, the use of Depth-First Search (DFS) in both the addTag and removeTag methods is a logical strategy for navigating the hierarchical structure of the tagged interval tree. DFS allows the algorithms to explore the tree and ensure that operations are applied correctly at different levels, respecting the parent-child relationships and the constraints imposed by the defined tree properties.¹ While standard interval tree searches for overlaps also involve tree traversal, often with optimizations based on the maximum high value stored in nodes, the specific application of DFS for insertion and deletion with the goal of maintaining a set of strict tree properties seems to be a characteristic tailored to the design of this particular tagged interval tree implementation.

Feature	Provided Implementation	node-interval-tree	@flatten-js/interval-tree	interval-tree-type
Underlying Data Structure	Likely a form of Binary Search Tree (not explicitly stated as self-balancing)	Augmented AVL Tree	Interval Binary Search Tree (based on Cormen et al.)	Red-Black Augmented Interval Tree
Add Interval/Tag	addTag with merging and tag property enforcement	insert (allows multiple entries for same interval)	insert (key-value pairs)	insert
Remove Interval/Tag	removeTag with complex overlap handling and re-hooking	remove (requires exact match of interval and data)	remove (requires exact match of key and value)	remove

Search/Query	Implicit through addTag, removeTag, hasTag, getFormattedText	search (overlap queries)	search, intersect_any	queryPoint, queryInterval, queryWithinInterval
Handles Overlapping Intervals	Enforces non-overlapping at the same level	Allows overlaps	Allows overlaps	Allows overlaps
Tagging Support	Explicitly designed for tagging with "Tag property"	Generic data associated with intervals	Generic value associated with interval key	Generic value associated with interval
Self-Balancing	Not explicitly mentioned	Yes (AVL)	Likely (based on BST and standard implementations)	Yes (Red-Black)

- **Conclusion: Synthesis of Similarities and Insights**

In summary, the provided tagged interval tree implementation shares the fundamental concept of using a tree structure to manage intervals with standard interval trees. It aims for efficient manipulation of text segments defined by start and end points. The operations for adding and removing tags implicitly involve searching and modifying the tree, which aligns with the general principles of interval tree algorithms.

However, this implementation distinguishes itself through several unique design choices tailored specifically for text tagging. The direct embedding of tags as a property within the IntervalNode and the strict enforcement of properties like Non-overlapping, Sorted, Merged, Inside, and Tag indicate a specific model for representing tagged text. The removeTag algorithm, with its intricate handling of various overlap scenarios and the re-hooking mechanism for managing child nodes, is particularly specialized.

The strict tree properties enforced by this implementation likely have significant implications for the functionality of a text editor using it. The non-overlapping and merged properties could simplify the rendering process by ensuring a clean and unambiguous structure of tagged text. The tag property might enforce a specific hierarchy of formatting or semantic meaning. However, the complexity of the

removeTag operation, while necessary for maintaining these properties, might have performance implications that differ from simpler interval removal operations in general-purpose libraries. A thorough evaluation of its performance characteristics would be essential for a practical application.

Ultimately, the provided tagged interval tree implementation appears to be a specialized solution designed with a particular model of tagged text in mind. While it draws upon the core principles of interval trees, its unique features and enforced properties suggest that it is more than just a general-purpose interval tree library. Its suitability would depend on the specific requirements of the text editing application, particularly regarding how tags should be structured, how overlaps should be handled, and the performance needs for adding and removing tags.

Works cited

1. Interval Tree | GeeksforGeeks, accessed April 24, 2025, <https://www.geeksforgeeks.org/interval-tree/>
2. Interval tree - Wikipedia, the free encyclopedia, accessed April 24, 2025, <https://ks3-cn-beijing.ksyun.com/attachment/c135186c1cc79c4d941f34443cc73f924>
3. Interval tree - Wikipedia, accessed April 24, 2025, https://en.wikipedia.org/wiki/Interval_tree
4. Practical Applications of Interval Tree - Stack Overflow, accessed April 24, 2025, <https://stackoverflow.com/questions/29653116/practical-applications-of-interval-tree>
5. Debian / python-intervaltree · GitLab, accessed April 24, 2025, <https://salsa.debian.org/debian/python-intervaltree>
6. FreshPorts -- devel/py-intervaltree: Editable interval tree data structure for Python 2 and 3, accessed April 24, 2025, <https://www.freshports.org/devel/py-intervaltree/>
7. python3-intervaltree : i386 : Oracular (24.10) : Ubuntu - Launchpad, accessed April 24, 2025, <https://launchpad.net/ubuntu/oracular/i386/python3-intervaltree>
8. chaimleib/intervaltree: A mutable, self-balancing interval tree. Queries may be by point, by range overlap, or by range containment. - GitHub, accessed April 24, 2025, <https://github.com/chaimleib/intervaltree>
9. intervaltree - PyPI, accessed April 24, 2025, <https://pypi.org/project/intervaltree/>
10. PyIntervalTree · PyPI, accessed April 24, 2025, <https://pypi.org/project/PyIntervalTree/>
11. konstantint/PyIntervalTree: A mutable, self-balancing interval tree. Queries may be by point, by range overlap, or by range containment. - GitHub, accessed April 24, 2025, <https://github.com/konstantint/PyIntervalTree>
12. node-interval-tree - NPM, accessed April 24, 2025, <https://www.npmjs.com/package/node-interval-tree>

13. alexbol99/flatten-interval-tree: Interval binary search tree - GitHub, accessed April 24, 2025, <https://github.com/alexbol99/flatten-interval-tree>
14. flatten-js/interval-tree - NPM, accessed April 24, 2025, <https://www.npmjs.com/package/@flatten-js/interval-tree>
15. IntervalTree - Documentation, accessed April 24, 2025, <https://alexbol99.github.io/flatten-interval-tree/IntervalTree.html>
16. pineapplemachine/interval-tree-type-js - GitHub, accessed April 24, 2025, <https://github.com/pineapplemachine/interval-tree-type-js>
17. circlingthesun/rb-interval-tree: Balanced interval tree for javascript - GitHub, accessed April 24, 2025, <https://github.com/circlingthesun/rb-interval-tree>
18. ucscXena/static-interval-tree: Fast overlapping interval queries in js. - GitHub, accessed April 24, 2025, <https://github.com/ucscXena/static-interval-tree>
19. An immutable interval tree in Javascript. - GitHub, accessed April 24, 2025, <https://github.com/davidisaaclee/interval-tree>
20. Segment tree implementation in JavaScript - chasin' clouds, accessed April 24, 2025, <http://www.chasinclouds.com/2012/02/segment-tree-implementation-in.html>
21. Difference Between Segment Trees, Interval Trees, Range Trees, and Binary Indexed Trees | Baeldung on Computer Science, accessed April 24, 2025, <https://www.baeldung.com/cs/tree-segment-interval-range-binary-indexed>
22. Bedtk: finding interval overlap with implicit interval tree - PMC, accessed April 24, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC8189672/>
23. Ask HN: Fast data structures for disjoint intervals? - Hacker News, accessed April 24, 2025, <https://news.ycombinator.com/item?id=41000844>
24. atkirtland/interval-tree-applet: Javascript visualization of Interval Trees - GitHub, accessed April 24, 2025, <https://github.com/atkirtland/interval-tree-applet>
25. Interval tree algorithm that supports merging of intervals with no overlap - Stack Overflow, accessed April 24, 2025, <https://stackoverflow.com/questions/2593774/interval-tree-algorithm-that-supports-merging-of-intervals-with-no-overlap>
26. Data structure for non overlapping ranges of integers? - Stack Overflow, accessed April 24, 2025, <https://stackoverflow.com/questions/19473671/data-structure-for-non-overlapping-ranges-of-integers>
27. A Web Browser and Editor - DSpace@MIT, accessed April 24, 2025, <https://dspace.mit.edu/bitstream/handle/1721.1/38137/35562181-MIT.pdf?sequence=2>
28. Text Editor Data Structures - invoke::thought(), accessed April 24, 2025, <https://cdacamar.github.io/data%20structures/algorithms/benchmarking/text%20editors/c++/editor-data-structures/>